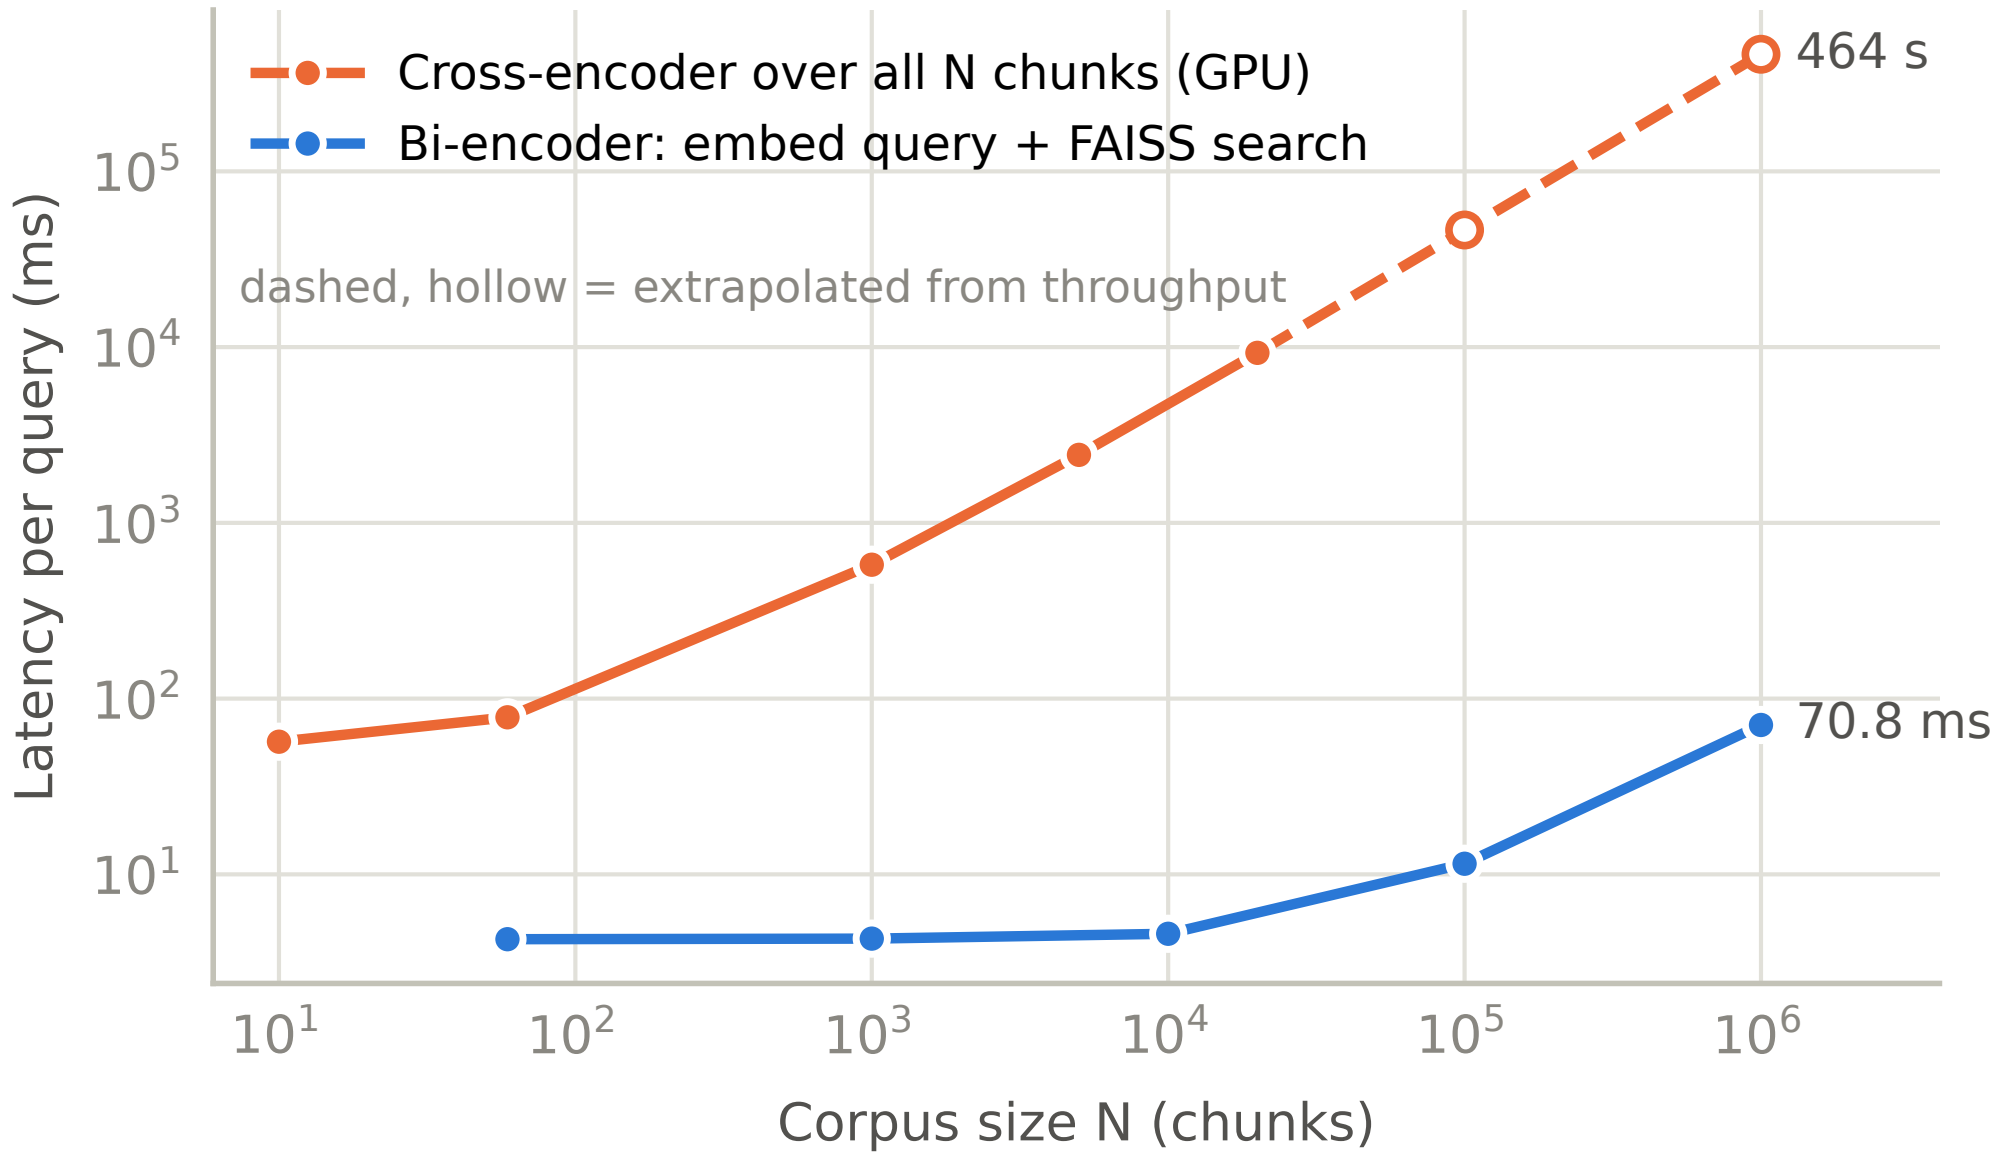
Cross-encoder over all N chunks (GPU)
Bi-encoder: embed query + FAISS search
dashed, hollow = extrapolated from throughput
464 s
70.8 ms
Latency per query (ms)
10^5
10^4
10^3
10^2
10^1
10^1
10^2
10^3
10^4
10^5
10^6
Corpus size N (chunks)